

The Sanctions Record

What courts have actually done about unverified AI output, 2023 to July 2026

An engineering read of 1,822 court decisions involving hallucinated AI content, and the controls that would have caught them.

Prepared by Firas Shaher, AI Solutions Architect. Data as published 1 August 2026. Report compiled 3 August 2026.

What this is

The AI Hallucination Cases database, maintained by Damien Charlotin, tracks court and tribunal decisions that address hallucinated AI content in filings. Its 1 August 2026 update holds 1,822 decisions from more than 40 jurisdictions. For this report I downloaded the full CSV export and counted the records myself rather than reusing a published summary. Every figure below is either a field in that record set or a number I verified in the cited decision or press report. Sources are listed on the last page.

Five numbers

- **Volume:** 16 decisions in 2023, 58 in 2024, 821 in 2025, and 913 in the first seven months of 2026.
- **Lawyers, not only self-represented filers:** 341 decisions involving lawyers in 2026 through July, already more than the 333 recorded for all of 2025.
- **Money:** 310 decisions record a monetary penalty. The 160 dollar-denominated sanctions above a nominal \$1 total about \$1.20 million, and the largest single judgment is \$110,204.
- **Careers:** 126 decisions carry a professional sanction flag such as a bar or disciplinary referral. The record is public, indexed and permanent, which is the part no fine covers.
- **Attribution:** In 1,603 of the 1,808 dated decisions, about 89 percent, no tool is named. Courts either inferred AI use or recorded the tool as unidentified.




Figure 1. Decisions per quarter, by decision date. Own count from the database CSV export, 1,808 records with a machine-readable date. Q3 2026 covers July only, and recent months undercount because decisions enter the database as they surface.

Who is filing hallucinated material

Self-represented litigants are the largest group in the database, 1,065 decisions against 710 for lawyers, with judges appearing 26 times and expert witnesses 13. Two things follow from that split, and they point in different directions.

- **The volume comes from unrepresented filers:** so the documents arriving from the other side now carry fabricated authority at a measurable rate. Verification is an inbound problem, not only an outbound one.
- **The professional exposure sits on the lawyer line:** which is the line growing fastest in absolute terms: 259 decisions in the second half of 2025, 307 in the first half of 2026.




Figure 2. Lawyers against self-represented litigants, by half-year. Own count from the CSV export. 2026 H2 covers July only.

What courts do about it

Warnings and admonishments are the most common outcome by a wide margin, so the median case is a bad day rather than a career event. The tail is where the cost lives, and the tail is not small: 364 decisions record money, 141 mention a referral to a bar or disciplinary body, 105 end in a claim, appeal or motion dismissed, and 80 have filings struck.




Figure 3. Outcome mix. My keyword classification of the free-text outcome field, so categories are not exclusive: a decision with both a fine and a referral counts in both rows. 359 records carry no outcome text at all, so treat these as counts of what the field states, not as a complete census of consequences.

The money

310 decisions record a monetary penalty. They span nine currencies, so no single total is honest without saying which slice it covers: 175 amounts in US dollars, 34 in shekels, 32 in Canadian dollars, 16 in euros, 8 in pounds, 7 in Australian dollars, 5 in Singapore dollars, and one each in Argentine pesos and Brazilian reais, plus 31 amounts recorded with no currency label. No conversion is applied anywhere in this report.

Inside the dollar slice, 160 sanctions above a nominal \$1 total \$1,201,678. The median dollar sanction was \$2,500 in 2025 and \$1,525 in 2026, which is the shape of the chart below: most orders are small and a handful carry the weight.




Figure 4. Size distribution of the 160 dollar-denominated sanctions above \$1. Own count from the CSV export.

Six decisions worth reading

Decision	Court and date	Consequence
Couvrette v. Wisnovsky	D. Oregon, fee judgment March 2026	\$110,204 against two lawyers, claims dismissed with prejudice. 15 fabricated citations and 8 fabricated quotations across three briefs filed over five months.
ByoPlanet International v. Johansson	S.D. Florida, 1 August 2025	Attorney ordered to pay \$85,567 of the defendants' fees and referred to the Florida Bar. ChatGPT named. Filings continued after the court put counsel on notice.
LiveVideo.AI Corp. v. Redstone	S.D. New York, 9 June 2026	\$80,056 adverse costs order plus a bar referral. Filings included leftover drafting prompts.
Lacey v. State Farm General Insurance	C.D. California, May 2025	\$31,100 against two law firms jointly. Nine defective citations out of 27, two to authorities that do not exist. The vendor, Thomson Reuters, said the errors did not originate in Westlaw or CoCounsel.
In re Rosslyn2016	S.D. Texas Bankruptcy, 14 July 2026	\$29,877 plus civil contempt and mandatory generative AI continuing legal education. Westlaw Precision named in the record and disputed by the vendor.
Scott v. Illinois Human Rights Commission	Illinois Appellate, 28 July 2026	\$15,000 fine plus a bar referral for a fabricated statutory quotation and a nonexistent appellate decision. Five days before this report was compiled.

The attribution gap

The database has a column for the AI tool involved. In 1,233 of the 1,808 dated decisions the tool is recorded as implied, meaning the court inferred AI use from the output itself. In another 370 it is recorded as unidentified. Only 205 decisions name a system, and 99 of those name ChatGPT. Nine records carry a vendor dispute, where the named vendor states publicly that the errors did not come from its product.




Figure 5. Tool attribution across the 1,808 dated decisions. Own count from the CSV export.

That gap is the part a builder should care about. Nobody can measure which system fails most often from this record, because the record mostly does not know. Attribution is a logging problem, not a legal one. If a document leaves your pipeline without a stored answer to "which model, which prompt, which retrieved source produced this citation", the provenance is gone by the time a judge asks.

The same failure, outside law

This is not a legal industry problem. On 30 July 2026 JFrog published an audit of 55 vulnerability advisories filed by one GitHub account: 54 could not be reproduced, disproved against the official SQLite tags with AddressSanitizer. One of them, CVE-2026-51302, still reads Critical 9.8, Unreviewed, no package listed, affected and patched versions unknown, with an exploit prediction score of 0.529 percent computed for a bug that does not exist. Record read 3 August 2026. Same shape as the court filings: plausible output entered a system of record with no reproduction step in front of it.

Controls that would have caught these

- **Bind citations to retrieval:** If no retrieval step returned the authority, the system does not render it. Generation proposes, retrieval authorises.
- **Check existence at egress:** Resolve every authority against an authoritative index when the document is assembled, not when the chat is happening. The gate belongs on the artifact that leaves.
- **Verify quotes as strings:** 493 decisions involve false quotes rather than invented cases, so existence checking alone passes them. Match quoted language against the source text.
- **Record provenance per claim:** Model, version, prompt and retrieved document IDs, stored immutably with the output. This is what turns "unidentified" into an answer.
- **Treat inbound documents as untrusted:** The largest group here is filers with no verification process at all. Check the authority in what you receive.
- **Route review by failure class, not by volume:** Every document that cites authority gets the citation gate. Sampling for tone catches nothing here.

None of this needs a new model. It needs a verification step with the power to block, and a log that survives the question.

Method and limits

- **Source of record:** CSV export of the AI Hallucination Cases database, published update 1 August 2026, downloaded and parsed for this report.
- **Coverage:** 1,808 of the 1,822 records carry both a machine-readable decision date and a recognised party value, so every time series here uses 1,808. Totals by party, jurisdiction and nature of hallucination are the database's own published counts.
- **Reporting lag:** Decisions enter the database as they surface, so the most recent months undercount. Read the right edge of every chart as provisional.
- **Scope:** The database tracks decisions that address hallucinated content, not every instance of AI misuse in filings, and it includes a small number of cases where AI use was alleged but not established.
- **Classification:** Outcome categories in Figure 3 are my keyword classification of a free-text field and are not mutually exclusive.

This is an engineering read of public records. It is not legal advice.

Sources

1. D. Charlotin, AI Hallucination Cases database, update of 1 August 2026. <https://www.damiencharlotin.com/hallucinations/> and its CSV export.
2. Couvrette v. Wisnovsky: Washington State Bar Association, NWSidebar, "Parade of Horribles: Federal Court in Oregon Surveys Sanctions for AI Fake Citations", 2 March 2026. <https://nwsidebar.wsba.org/2026/03/02/parade-of-horribles-federal-court-in-oregon-surveys-sanctions-for-ai-fake-citations/>
3. ByoPlanet International LLC v. Johansson: sanctions order, S.D. Fla., 1 August 2025, via FindLaw. <https://caselaw.findlaw.com/court/us-dis-crt-sd-fla/117513089.html>
4. Lacey v. State Farm General Insurance Co., 2025 WL 1363069 (C.D. Cal.): Los Angeles County Bar Association, "Artificial Intelligence Goes to Court". <https://lacba.org>
5. In re Rosslyn2016, LLC, S.D. Tex. Bankruptcy, 14 July 2026: decision and vendor-dispute note in the database record. <https://www.damiencharlotin.com/hallucinations/>
6. Scott v. Illinois Human Rights Commission: E. Volokh, Reason, 31 July 2026. <https://reason.com/volokh/2026/07/31/>
7. M. Hiltzik, "AI hallucinations are a growing problem for the legal profession", Los Angeles Times, 22 May 2025. <https://www.latimes.com/business/story/2025-05-22/ai-lawyers-hallucinations-a-growing-problem-for-lawyers>
8. J. Koebler and J. Roscoe, "18 Lawyers Caught Using AI Explain Why They Did It", 404 Media, 30 September 2025. <https://www.404media.co/18-lawyers-caught-using-ai-explain-why-they-did-it/>
9. JFrog Security Research, "SQLite critical CVEs or LLM slops?", 30 July 2026. <https://research.jfrog.com/post/sqlite-critical-cves-or-llm-slops/>
10. CVE-2026-51302, GitHub Advisory Database, GHSA-vrg3-8p22-cwh8, record read 3 August 2026. <https://github.com/advisories/GHSA-vrg3-8p22-cwh8>